

CryptoWhiz

Submitted by:
Yash Karan

BE Third Year
CSE/COE

Submitted to:
Dr. Anjula Mehto
Assistant Professor



Computer Science and Engineering Department
Thapar Institute of Engineering and Technology, Patiala

November 2025

TABLE OF CONTENTS

S. No	Topic	Page No.
1	Introduction	3
2	Problem Statement	4
3	Overview of the Dataset used	5
4	Project workflow	7
5	Results	17
6	Conclusion	20

Introduction

CryptoWhiz represents a cutting-edge exploration into the intersection of artificial intelligence, machine learning, and financial technology, specifically designed to navigate the complex and volatile world of cryptocurrency trading. This sophisticated educational project demonstrates how multiple AI agents can work in harmony to analyze market conditions, assess risks, and generate intelligent trading decisions for a portfolio of the top 10 cryptocurrencies. At its core, CryptoWhiz employs a multi-agent architecture where specialized components work together like a well-coordinated team: a Technical Analysis Agent utilizing Long Short-Term Memory (LSTM) neural networks interprets price charts and predicts market movements; a Sentiment Analysis Agent powered by FinBERT processes news headlines and social media to gauge market psychology; a Risk Management Agent leverages GARCH volatility models to prevent catastrophic losses and optimize position sizing; and finally, a Master Trading Agent employing Deep Reinforcement Learning synthesizes all these signals to make the ultimate buy, sell, or hold decisions. Built entirely with free, open-source tools and publicly available data sources including the CoinGecko API and Yahoo Finance, CryptoWhiz serves as both a comprehensive learning resource for students and enthusiasts interested in algorithmic trading and a practical demonstration of how modern AI techniques can be applied to real-world financial challenges. While the system showcases the potential of AI-driven trading strategies through features like automated technical indicator calculation, sentiment scoring, volatility forecasting, and intelligent position management, it is crucial to emphasize that this project is purely educational in nature and should never be used for actual trading without extensive modifications, professional consultation, and a thorough understanding of the substantial risks involved in cryptocurrency markets. Whether you're a data science student looking to understand quantitative finance, a developer interested in building AI trading systems, or simply curious about how machine learning can be applied to financial markets, CryptoWhiz offers a hands-on, practical introduction to the fascinating world of algorithmic cryptocurrency trading while maintaining a strong focus on responsible development and risk awareness.

Problem Statement

Cryptocurrency markets exhibit volatility levels that far exceed traditional financial markets, creating both opportunities and dangers that many traders are ill-equipped to navigate. Daily price fluctuations of 5-10% are routine for major cryptocurrencies like Bitcoin and Ethereum, while smaller-cap altcoins can experience 20-50% swings in single trading sessions. Flash crashes occur with disturbing regularity where prices plummet 30-40% in minutes due to cascading liquidations, exchange glitches, or large market orders, often followed by equally rapid recoveries that trap traders who sold in panic. Black swan events such as exchange hacks, regulatory announcements, macroeconomic shocks, or influential figures' social media posts can trigger unprecedented market moves that exceed historical volatility patterns and invalidate risk models based on past data. Leverage amplifies both gains and losses, with many cryptocurrency exchanges offering 50x, 100x, or even 125x leverage that can liquidate entire positions from minor price movements. Liquidation cascades create positive feedback loops where margin calls trigger automatic selling that drives prices lower, triggering more margin calls in a devastating spiral. The lack of circuit breakers and trading halts in cryptocurrency markets means that crashes continue unabated without the cooling-off periods that traditional markets employ to prevent panic. Most retail traders lack formal training in risk management, position sizing, portfolio diversification, and the quantitative techniques necessary to protect capital in such extreme conditions, leading to devastating losses that wipe out months or years of gains in single catastrophic events.

Overview of the Dataset used

CryptoWhiz utilizes a combination of free, publicly available data sources to perform comprehensive cryptocurrency market analysis. The primary data source is the CoinGecko API, which provides daily historical price data for cryptocurrencies including closing prices, 24-hour trading volumes, and approximated OHLCV (Open, High, Low, Close, Volume) values. This API is accessed through the endpoint `/api/v3/coins/{id}/market_chart` and requires no authentication or API key, making it freely accessible for educational purposes. The system retrieves data for configurable time periods ranging from one month to one year, with the default setting at six months (180 days). Each data point includes a timestamp, closing price in USD, trading volume, and derived open, high, and low prices. As a fallback mechanism, the system also integrates with Yahoo Finance through the `yfinance` Python library, which provides more accurate OHLCV data when available and serves as a backup when CoinGecko experiences rate limiting or temporary outages.

The cryptocurrency portfolio analyzed by CryptoWhiz consists of the top 10 cryptocurrencies by market capitalization, including Bitcoin (BTC-USD), Ethereum (ETH-USD), Binance Coin (BNB-USD), Ripple (XRP-USD), Cardano (ADA-USD), Dogecoin (DOGE-USD), Solana (SOL-USD), Tron (TRX-USD), Polkadot (DOT-USD), and Tether (USDT-USD). These cryptocurrencies were selected based on their high liquidity, significant trading volumes, availability of sufficient historical data spanning at least two years, and diverse representation of different cryptocurrency categories including Layer 1 blockchains, DeFi tokens, and meme coins. For each cryptocurrency, the system retrieves approximately 180 data points for a six-month analysis period, resulting in a comprehensive dataset of around 19,800 total data points across all features and cryptocurrencies.

From the raw price data, CryptoWhiz calculates several technical indicators that are widely used in financial analysis. The Relative Strength Index (RSI) is computed using a 14-day period, measuring momentum on a scale from 0 to 100, where values above 70 indicate overbought conditions and values below 30 indicate oversold conditions. Simple Moving Averages (SMA) are calculated for both 20-day and 50-day periods to identify short-term and medium-term trends respectively. The Moving Average Convergence Divergence (MACD) indicator is derived by calculating the difference between 12-day and 26-day exponential moving averages, along with a 9-day signal line and histogram. Daily returns are calculated as logarithmic returns to measure percentage price changes, and historical volatility is computed as the standard deviation of these returns over a rolling 30-day window. All of these indicators are stored in a Pandas DataFrame structure with a DatetimeIndex, allowing for efficient time-series analysis and manipulation.

For sentiment analysis, CryptoWhiz employs FinBERT, a specialized Natural Language Processing model from the HuggingFace Model Hub (model ID: ProsusAI/finbert). FinBERT is a variant of BERT (Bidirectional Encoder Representations from Transformers) that has been specifically fine-tuned on financial domain text including news articles, analyst reports, and earnings calls. The model processes cryptocurrency-related news

headlines and market commentary to generate sentiment scores across three categories: positive, negative, and neutral, each with associated probability values ranging from 0 to 1. A compound sentiment score is calculated by subtracting the negative probability

from the positive probability, resulting in a final score ranging from -1 (very negative) to +1 (very positive). In the current educational implementation, the system processes mock news headlines to demonstrate the sentiment analysis capability, though the architecture is designed to easily integrate with real-time news APIs such as NewsAPI, CryptoPanic, or social media platforms for production use. If FinBERT fails to load due to system constraints, the system includes a fallback rule-based sentiment analyzer that uses keyword matching to identify bullish and bearish terms in text.

The data preprocessing pipeline includes normalization techniques specifically designed for machine learning applications. For LSTM neural network training, price data is normalized using MinMaxScaler to transform values into a 0-1 range, which significantly improves neural network performance and training stability. Technical indicators like RSI are kept in their original ranges since they are already bounded and use standard interpretation thresholds. The dataset is split into training and testing sets using an 80-20 split, where the oldest 80% of data is used for training the LSTM model and the most recent 20% is reserved for testing, ensuring no data leakage occurs where future information could improperly influence predictions about the past. Missing values in the time series are handled through forward-filling for gaps up to three days, while longer gaps result in data exclusion to maintain quality. The system also implements outlier detection to flag any single-day price changes exceeding 50% for manual review, and filters out days with zero or suspiciously low trading volume to ensure data quality and reliability throughout the analysis pipeline.

Project Workflow

The CryptoWhiz trading system operates through a sophisticated multi-stage pipeline where data flows through specialized AI agents, each performing specific analytical tasks before converging at the Master Trading Agent for final decision-making. This document explains each stage of the workflow with corresponding code examples to illustrate the process.

Stage 1: System Initialization

The workflow begins with initializing all components of the CryptoWhiz system. During initialization, the system sets up five key modules that will work together throughout the analysis process.

```
# Initialize the complete CryptoWhiz system
cryptowhiz = CryptoWhiz()
```

During this initialization phase, the Technical Analysis Agent prepares its LSTM neural network architecture, the Sentiment Analysis Agent loads the pre-trained FinBERT model from HuggingFace with approximately 110 million parameters, the Risk Management Agent configures GARCH volatility modeling parameters, and the Master Trading Agent builds its Deep Q-Network with three possible actions (buy, sell, hold). The system also initializes empty portfolio and trade history tracking structures that will be populated during analysis.

Stage 2: Data Acquisition

Once the system is initialized, the data acquisition module begins fetching historical price data for all ten cryptocurrencies from external APIs. This stage implements intelligent rate limiting and fallback mechanisms to ensure reliable data retrieval.

The data acquisition process iterates through each cryptocurrency symbol, first attempting to retrieve data from the CoinGecko API using the endpoint /api/v3/coins/{id}/market_chart. If CoinGecko fails or is rate-limited, the system automatically falls back to Yahoo Finance through the yfinance library. Each successful retrieval produces a Pandas DataFrame containing approximately 180 rows of daily OHLCV data spanning six months. The system implements a 1.5-second delay between API calls to respect rate limits and avoid being blocked. For each cryptocurrency, the returned DataFrame includes columns for Open, High, Low, Close prices, Volume, and a

Symbol identifier, all indexed by timestamp for efficient time-series operations.

```
def fetch_from_coingecko(self, symbol, days=180):
    """Fetch data from CoinGecko API (free, no auth needed)"""
    try:
        coin_id = self.coingecko_map.get(symbol)
        if not coin_id:
            return None

        url = f"https://api.coingecko.com/api/v3/coins/{coin_id}/market_chart"
        params = {
            'vs_currency': 'usd',
            'days': days,
            'interval': 'daily'
        }

        response = requests.get(url, params=params, timeout=10)

        if response.status_code == 200:
            data = response.json()

            # Convert to DataFrame
            prices = data['prices']
            volumes = data['total_volumes']

            df = pd.DataFrame(prices, columns=['timestamp', 'Close'])
            df['Volume'] = [v[1] for v in volumes]
            df['timestamp'] = pd.to_datetime(df['timestamp'], unit='ms')
            df.set_index('timestamp', inplace=True)

            # Add OHLC (approximate from close prices)
            df['Open'] = df['Close'].shift(1).fillna(df['Close'])
            df['High'] = df['Close'] * 1.02 # Approximate
            df['Low'] = df['Close'] * 0.98 # Approximate
            df['Symbol'] = symbol

            return df
        else:
            return None

    except Exception as e:
        print(f" ⚠️ CoinGecko error for {symbol}: {e}")
        return None

def fetch_price_data(self, symbol, period='1y', interval='1d'):
    """Fetch historical price data with fallback sources"""
    # Try Yahoo Finance first
    try:
        ticker = yf.Ticker(symbol)
        df = ticker.history(period=period, interval=interval)
        if not df.empty:
            df['Symbol'] = symbol
            return df
    except Exception as e:
        pass # Silent fail, try alternative

    # Fallback to CoinGecko
    days_map = {'1mo': 30, '3mo': 90, '6mo': 180, '1y': 365}
    days = days_map.get(period, 180)

    return self.fetch_from_coingecko(symbol, days=days)
```

```

def fetch_all_cryptos(self, period='6mo'):
    """Fetch data for all top 10 cryptocurrencies with rate limiting"""
    data = {}
    print("Fetching cryptocurrency data...")
    print("(Using CoinGecko API )\n")

    for i, crypto in enumerate(self.top_10_cryptos, 1):
        print(f"[{i}/{len(self.top_10_cryptos)}] Downloading {crypto}...", end=' ')
        df = self.fetch_price_data(crypto, period)

        if df is not None and not df.empty:
            data[crypto] = df
            print(f"\u2713 ({len(df)} data points)")
        else:
            print("X Failed")

        # Rate Limiting: wait between requests
        if i < len(self.top_10_cryptos):
            time.sleep(1.5) # 1.5 second delay to respect API Limits

    print("\u2713 Successfully fetched {len(data)}/{len(self.top_10_cryptos)} cryptocurrencies\n")
    return data

```

Stage 3: Technical Indicator Calculation

After successfully acquiring raw price data, the system calculates a comprehensive set of technical indicators that provide insights into price trends, momentum, and potential reversal points. These indicators form the foundation for the LSTM model's predictions.

The RSI calculation begins by computing the daily price differences, separating them into gains and losses, and calculating 14-day rolling averages of each. The relative strength (RS) ratio of average gains to average losses is then transformed into the RSI value bounded between 0 and 100. Moving averages are computed using rolling windows of 20 and 50 days, smoothing out short-term price fluctuations to reveal underlying trends. The MACD indicator subtracts a 26-day exponential moving average from a 12-day exponential moving average, with the resulting MACD line's 9-day exponential moving average serving as the signal line. These indicators are appended as new columns to the existing DataFrame, enriching the dataset with technical analysis features that capture different aspects of market behavior such as momentum, trend direction, and potential overbought or oversold conditions.

```

def calculate_indicators(self, df):
    """Calculate technical indicators"""

    # RSI
    delta = df['Close'].diff()
    gain = (delta.where(delta > 0, 0)).rolling(window=14).mean()
    loss = (-delta.where(delta < 0, 0)).rolling(window=14).mean()
    rs = gain / loss
    df['RSI'] = 100 - (100 / (1 + rs))

    # Moving Averages
    df['SMA_20'] = df['Close'].rolling(window=20).mean()
    df['SMA_50'] = df['Close'].rolling(window=50).mean()

    # MACD
    exp1 = df['Close'].ewm(span=12, adjust=False).mean()
    exp2 = df['Close'].ewm(span=26, adjust=False).mean()
    df['MACD'] = exp1 - exp2
    df['Signal_Line'] = df['MACD'].ewm(span=9, adjust=False).mean()

    return df

```

Stage 4: LSTM Model Training

With technical indicators calculated, the Technical Analysis Agent prepares the data for deep learning and trains an LSTM neural network to predict future price movements based on historical patterns. This stage involves data preprocessing, sequence creation, and neural network training.

```
class TechnicalAnalysisAgent:
    """LSTM-based price prediction and technical analysis"""

    def __init__(self, lookback=60):
        self.lookback = lookback
        self.scaler = MinMaxScaler()
        self.model = None

    def prepare_data(self, df):
        """Prepare data for LSTM training"""
        # Use Close price
        data = df['Close'].values.reshape(-1, 1)
        scaled_data = self.scaler.fit_transform(data)

        X, y = [], []
        for i in range(self.lookback, len(scaled_data)):
            X.append(scaled_data[i-self.lookback:i, 0])
            y.append(scaled_data[i, 0])

        X, y = np.array(X), np.array(y)
        X = np.reshape(X, (X.shape[0], X.shape[1], 1))

        return X, y, scaled_data
    def train(self, df, epochs=10, batch_size=32):
        """Train LSTM model"""
        X, y, _ = self.prepare_data(df)

        # Split data
        split = int(0.8 * len(X))
        X_train, X_test = X[:split], X[split:]
        y_train, y_test = y[:split], y[split:]

        self.model = self.build_model()

        print(f"Training LSTM model with {len(X_train)} samples...")
        self.model.fit(
            X_train, y_train,
            batch_size=batch_size,
            epochs=epochs,
            validation_data=(X_test, y_test),
            verbose=0
        )

    return self.model
```

The data preparation phase extracts closing prices and normalizes them to a 0-1 range using MinMaxScaler, which is crucial for LSTM convergence and performance. The system then creates sequential input data using a 60-day lookback window, where each training example consists of 60 consecutive days of normalized prices as input (X) and the next day's price as the target output (y). These sequences are reshaped into three-

dimensional arrays with dimensions (samples, timesteps, features) as required by LSTM layers. The dataset is split 80-20 into training and testing sets chronologically to prevent data leakage. The LSTM architecture consists of two stacked LSTM layers with 50 units each, dropout layers with 20% probability to prevent overfitting, and dense layers that compress the learned representations into a single price prediction. The model is compiled with the Adam optimizer and mean squared error loss function, then trained for five epochs with a batch size of 32, learning to recognize temporal patterns and relationships in the sequential price data.

Stage 5: Price Prediction

After training, the LSTM model generates predictions for the next day's price movement. The Technical Analysis Agent uses the most recent 60 days of price data to forecast tomorrow's price and determines whether this suggests a buying or selling opportunity.

```
def predict_next(self, df):
    """Predict next price movement"""
    if self.model is None:
        return None

    _, _, scaled_data = self.prepare_data(df)
    last_sequence = scaled_data[-self.lookback:].reshape(1, self.lookback, 1)

    prediction = self.model.predict(last_sequence, verbose=0)
    prediction = self.scaler.inverse_transform(prediction)

    current_price = df['Close'].iloc[-1]
    predicted_price = prediction[0][0]

    return {
        'current': current_price,
        'predicted': predicted_price,
        'change_pct': ((predicted_price - current_price) / current_price) * 100,
        'signal': 'BUY' if predicted_price > current_price else 'SELL'
    }
```

The prediction process begins by extracting the most recent 60-day sequence of normalized prices from the dataset, which serves as the input context for the LSTM model. This sequence is reshaped into the required three-dimensional format and fed through the trained neural network, which processes the temporal patterns through its LSTM layers and produces a single normalized value representing tomorrow's predicted price. The prediction is then inverse-transformed back to the original price scale using the same scaler that was applied during training. The system calculates the percentage change between the current price and predicted price, providing a quantitative measure of expected movement. Based on this comparison, a simple trading signal is generated: if the predicted price exceeds the current price, the signal is 'BUY', indicating an upward trend; otherwise, it's 'SELL', suggesting a downward trend. The function returns a dictionary containing the current price, predicted price, expected percentage change, and the trading signal, which will be used by the Master Trading Agent in its final decision-making process.

Stage 6: Sentiment Analysis

Parallel to technical analysis, the Sentiment Analysis Agent processes textual information about the cryptocurrency to gauge market psychology and investor mood. This provides a qualitative dimension to complement the quantitative technical indicators.

```
def get_market_sentiment(self, symbol, news_list=None):
    """
    Get overall market sentiment for a cryptocurrency
    In production, fetch real news using NewsAPI or CryptoPanic
    """

    if news_list is None:
        # Mock news headlines
        news_list = [
            f"{symbol} shows strong momentum with institutional buying",
            f"Market analysts bullish on {symbol} future prospects",
            f"{symbol} faces regulatory concerns in key markets"
        ]

    sentiments = []
    for news in news_list:
        sent = self.analyze_text(news)
        sentiments.append(sent['compound'])

    avg_sentiment = np.mean(sentiments) if sentiments else 0

    return {
        'sentiment_score': avg_sentiment,
        'signal': 'POSITIVE' if avg_sentiment > 0.1 else 'NEGATIVE' if avg_sentiment < -0.1 else 'NEUTRAL',
        'confidence': abs(avg_sentiment)
    }
```

The sentiment analysis process begins by collecting relevant news headlines or text snippets about the cryptocurrency. In the current implementation, mock headlines are generated for demonstration, but the architecture supports integration with real-time news APIs. Each text snippet is tokenized using FinBERT's specialized tokenizer, which converts the text into numerical token IDs that the model can process, with truncation at 512 tokens and appropriate padding. The tokenized input is passed through the FinBERT transformer model, which processes the text through 12 layers of attention mechanisms to understand context and meaning. The model outputs logits for three sentiment classes (negative, neutral, positive), which are converted to probabilities using the softmax function. A compound sentiment score is calculated by subtracting the negative probability from the positive probability, yielding a value between -1 (very negative) and +1 (very positive). This process is repeated for each news headline, and the individual compound scores are averaged to produce an overall sentiment assessment for the cryptocurrency. The final sentiment signal is classified as POSITIVE if the average exceeds 0.1, NEGATIVE if below -0.1, or NEUTRAL otherwise, with confidence measured by the absolute magnitude of the score.

Stage 7: Risk Assessment

The Risk Management Agent performs a comprehensive volatility analysis using GARCH modeling to quantify the risk associated with trading the cryptocurrency and determine appropriate position sizing to protect against excessive losses.

```
def assess_risk(self, df):
    """Comprehensive risk assessment"""
    returns = self.calculate_returns(df)

    if len(returns) < 50:
        return None

    # Fit GARCH model
    garch_fit = self.fit_garch(returns)

    # Calculate volatility
    volatility = returns.std()

    # Calculate VaR
    var_95 = self.calculate_var(returns, 0.95)

    # Forecast volatility
    forecast_vol = volatility
    if garch_fit is not None:
        forecast = garch_fit.forecast(horizon=1)
        forecast_vol = np.sqrt(forecast.variance.values[-1, :][0])

    # Risk assessment
    risk_level = 'LOW'
    if forecast_vol > 5:
        risk_level = 'HIGH'
    elif forecast_vol > 3:
        risk_level = 'MEDIUM'

    # Position sizing
    position_size = self.max_position_size * (1 - min(forecast_vol / 10, 0.8))

    return {
        'volatility': volatility,
        'forecast_volatility': forecast_vol,
        'var_95': var_95,
        'risk_level': risk_level,
        'recommended_position': position_size,
        'sharpe_ratio': returns.mean() / volatility if volatility > 0 else 0
    }
```

Risk assessment begins by computing logarithmic returns, which represent the continuously compounded rate of change in prices and are preferred over simple returns due to their mathematical properties in time-series analysis. The system fits a GARCH(1,1) model to these returns, which captures volatility clustering, the tendency for periods of high volatility to be followed by high volatility and low volatility by low volatility. The GARCH model estimates both the current volatility level and forecasts tomorrow's expected volatility by modeling how today's volatility depends on yesterday's volatility and the magnitude of recent price shocks. Historical volatility is calculated as the standard deviation of returns, providing a baseline measure of price variability. Value at Risk (VaR)

at 95% confidence is computed as the 5th percentile of the return distribution, representing the maximum expected loss on 19 out of 20 days. The system classifies risk level as HIGH when forecast volatility exceeds 5% daily, MEDIUM between 3-5%, and LOW below 3%. Position sizing is dynamically adjusted based on forecast volatility, starting from a maximum position size of 20% and reducing it as volatility increases, ensuring smaller positions during turbulent periods to limit potential losses. The Sharpe ratio is calculated by dividing average returns by volatility, measuring risk-adjusted performance where higher values indicate better return per unit of risk taken. All these metrics are packaged into a comprehensive risk profile that guides the final trading decision.

Stage 8: Master Decision Making

The Master Trading Agent serves as the final decision-maker, integrating signals from technical analysis, sentiment analysis, and risk management into a unified trading decision using a Deep Q-Network that has learned optimal action selection through reinforcement learning principles.

```
def make_decision(self, ta_signal, sentiment_signal, risk_assessment):
    """
    Integrate all signals to make final trading decision

    Returns: 0 (SELL), 1 (HOLD), 2 (BUY)
    """

    # Create state vector
    state = np.array([
        1 if ta_signal['signal'] == 'BUY' else 0,
        ta_signal['change_pct'],
        sentiment_signal['sentiment_score'],
        risk_assessment['volatility'],
        risk_assessment['forecast_volatility'],
        risk_assessment['var_95'],
        risk_assessment['sharpe_ratio'],
        1 if risk_assessment['risk_level'] == 'HIGH' else 0,
        risk_assessment['recommended_position'],
        sentiment_signal['confidence']
    ]).reshape(1, -1)

    # Epsilon-greedy action selection
    if np.random.rand() <= self.epsilon:
        action = np.random.randint(self.action_size)
    else:
        q_values = self.model.predict(state, verbose=0)
        action = np.argmax(q_values[0])

    # Apply risk override
    if risk_assessment['risk_level'] == 'HIGH':
        action = min(action, 1) # Only HOLD or SELL in high risk

    action_map = {0: 'SELL', 1: 'HOLD', 2: 'BUY'}

    return {
        'action': action_map[action],
        'confidence': np.max(self.model.predict(state, verbose=0)),
        'position_size': risk_assessment['recommended_position'],
        'reasoning': self._explain_decision(ta_signal, sentiment_signal, risk_assessment, action_map[action])
    }
```

Stage 9: Results Compilation and Display

After analyzing each cryptocurrency individually, the system compiles all results and generates portfolio-level recommendations, ranking opportunities by their attractiveness based on confidence levels and risk-adjusted potential returns.

```
def _display_analysis(self, symbol, ta, sentiment, risk, decision):
    """Display comprehensive analysis results"""
    print(f"\n📊 TECHNICAL ANALYSIS")
    print(f" Current Price: ${ta['current']:.2f}")
    print(f" Predicted Price: ${ta['predicted']:.2f}")
    print(f" Expected Change: {ta['change_pct']:+.2f}%")
    print(f" Signal: {ta['signal']}")

    print(f"\n🟡 SENTIMENT ANALYSIS")
    print(f" Sentiment Score: {sentiment['sentiment_score']:+.3f}")
    print(f" Signal: {sentiment['signal']}")
    print(f" Confidence: {sentiment['confidence']:.2%}")

    print(f"\n⚠️ RISK ASSESSMENT")
    print(f" Current Volatility: {risk['volatility']:.2f}%")
    print(f" Forecast Volatility: {risk['forecast_volatility']:.2f}%")
    print(f" VaR (95%): {risk['var_95']:.2f}%")
    print(f" Risk Level: {risk['risk_level']}")
    print(f" Sharpe Ratio: {risk['sharpe_ratio']:.3f}")

    print(f"\n🔴 MASTER DECISION")
    print(f" Action: {decision['action']}")
    print(f" Confidence: {decision['confidence']:.2%}")
    print(f" Position Size: {decision['position_size']:.1%}")
    print(f" Reasoning: {decision['reasoning']}")

def _generate_portfolio_recommendations(self, results):
    """Generate portfolio allocation recommendations"""
    print("\n" + "="*60)
    print("🔗 PORTFOLIO RECOMMENDATIONS")
    print("="*60)

    buy_signals = []
    for symbol, analysis in results.items():
        if analysis['decision']['action'] == 'BUY':
            buy_signals.append({
                'symbol': symbol.replace('-USD', ''),
                'position_size': analysis['decision']['position_size'],
                'confidence': analysis['decision']['confidence']
            })

    if buy_signals:
        # Sort by confidence
        buy_signals.sort(key=lambda x: x['confidence'], reverse=True)

        print("\n🟡 TOP BUY RECOMMENDATIONS:")
        for i, signal in enumerate(buy_signals[:5], 1):
            print(f" {i}. {signal['symbol']}: {signal['position_size']:.1%} allocation (confidence: {signal['confidence']:.2%})")
    else:
        print("\n🚫 No strong buy signals detected. Consider holding current positions.")

    print("\n" + "="*60)
```

The results display phase presents a structured summary of the complete analysis for each cryptocurrency, organized into four main sections. The technical analysis section shows the current market price, the LSTM-predicted price for tomorrow, and the expected percentage change, providing a clear view of anticipated price movement. The sentiment analysis section displays the compound sentiment score on a -1 to +1 scale and classifies it as POSITIVE, NEGATIVE, or NEUTRAL, revealing market psychology. The risk assessment section presents key metrics including historical volatility, risk classification, and the Sharpe ratio for risk-adjusted performance evaluation. The master decision section announces the final trading action (BUY, SELL, or HOLD), the recommended position size as a percentage of portfolio allocation, and a synthesized reasoning statement that explains the decision logic. After analyzing all ten cryptocurrencies, the system generates portfolio-level recommendations by filtering for BUY signals, sorting them by confidence level, and presenting the top five opportunities with their recommended allocation percentages. This prioritized list helps users understand which cryptocurrencies present the most attractive risk-reward profiles based on the AI's comprehensive analysis, enabling informed portfolio construction decisions while maintaining appropriate risk diversification across multiple assets.

Results

Result Category 1: Technical Analysis Predictions

The Technical Analysis Agent produces price forecasts and trend signals by processing historical price data through trained LSTM neural networks that have learned to recognize patterns in cryptocurrency market movements.

Current Price: \$43,250.00

Predicted Price: \$44,180.00

Expected Change: +2.15%

Signal: BUY

Result Category 2: Sentiment Analysis Scores

The Sentiment Analysis Agent processes cryptocurrency-related news and market commentary through the FinBERT transformer model to quantify market psychology and investor mood surrounding each asset.

Sentiment Score: +0.156

Signal: POSITIVE

Confidence: 15.6%

Sentiment Breakdown:

- Positive Probability: 0.45 (45% chance news is positive)
- Neutral Probability: 0.35 (35% chance news is neutral)
- Negative Probability: 0.20 (20% chance news is negative)
- Compound Score: +0.25 (calculated as positive - negative)

Result Category 3: Risk Assessment Metrics

The Risk Management Agent employs GARCH volatility modeling to quantify the risk profile of each cryptocurrency and determine appropriate position sizing to protect against excessive losses.

Current Volatility: 3.24%

Forecast Volatility: 3.87%

Value at Risk (95%): -4.52%

Risk Level: MEDIUM

Recommended Position Size: 15.2%

Sharpe Ratio: 0.342

Volatility Measurements:

- Historical Volatility (3.24%): The standard deviation of daily returns over the past six months, indicating that on a typical day, the price moves up or down by approximately 3.24% from its average. Lower volatility indicates more stable, predictable price movements, while higher volatility suggests erratic and unpredictable behavior.
- Forecast Volatility (3.87%): The GARCH model's prediction for tomorrow's expected volatility based on recent volatility patterns and the magnitude of recent price shocks. This forward-looking metric helps anticipate periods of increased market turbulence. The forecast being higher than historical volatility suggests increasing market uncertainty.

Value at Risk (VaR):

- VaR at 95% Confidence (-4.52%): Statistical measure indicating that on 19 out of 20 trading days, losses should not exceed 4.52% of the position value. This metric provides a quantitative estimate of potential losses under normal market conditions. For example, a \$10,000 position would not be expected to lose more than \$452 on 95% of days.

Risk Level Classification:

- LOW: Forecast volatility below 3% daily (stable, low-risk assets)
- MEDIUM: Forecast volatility between 3-5% daily (moderate risk)
- HIGH: Forecast volatility above 5% daily (highly volatile, risky assets)

Sharpe Ratio (0.342):

- Measures risk-adjusted returns by dividing average returns by volatility. Higher values indicate better return per unit of risk taken. A Sharpe ratio of 0.342 suggests modest risk-adjusted performance, meaning the returns barely compensate for the volatility risk. Values above 1.0 are generally considered good, above 2.0 excellent, and below 0 indicate losses.

Result Category 4: Master Trading Decisions

The Master Trading Agent synthesizes all signals from technical analysis, sentiment analysis, and risk management through its Deep Q-Network to produce final trading recommendations with supporting reasoning.

Action: BUY

Confidence: 68.3%

Position Size: 15.2%

Reasoning: Technical analysis predicts +2.15% price increase | Market sentiment is POSITIVE

Possible Actions:

1. BUY: Enter or increase position in the cryptocurrency
2. HOLD: Maintain current position without changes
3. SELL: Exit or reduce position in the cryptocurrency

Decision Confidence:

- Represents the Deep Q-Network's certainty in its recommendation based on the expected reward (Q-value) for the chosen action
- Higher confidence (above 70%) suggests strong alignment among all signals
- Lower confidence (below 50%) indicates conflicting signals or uncertainty

Result Category 5: Portfolio-Level Recommendations

After analyzing all ten cryptocurrencies individually, the system generates consolidated portfolio recommendations that rank opportunities by their attractiveness and provide allocation guidance.

🔥 TOP BUY RECOMMENDATIONS:

1. BTC: 18.5% allocation (confidence: 85.2%)
2. ETH: 16.8% allocation (confidence: 78.4%)
3. SOL: 14.2% allocation (confidence: 71.9%)
4. DOT: 12.6% allocation (confidence: 66.3%)
5. ADA: 11.1% allocation (confidence: 63.7%)

Conclusion

The CryptoWhiz project represents a comprehensive exploration of how modern artificial intelligence techniques can be applied to cryptocurrency market analysis, demonstrating both the exciting potential and important limitations of autonomous trading systems. By successfully integrating LSTM neural networks, transformer-based sentiment analysis, GARCH volatility modeling, and Deep Q-Networks into a cohesive multi-agent architecture, the project creates a system that provides genuine analytical value while maintaining appropriate caution about its educational nature and inherent constraints. The development process required mastery of diverse technical skills spanning machine learning, software engineering, quantitative finance, and data science, resulting in a portfolio-worthy project that showcases interdisciplinary problem-solving abilities highly valued in today's technology-driven financial industry.

Perhaps most importantly, the project emphasizes responsible AI development through its strong disclaimers, risk management focus, and educational framing. In an era where AI capabilities are advancing rapidly and becoming increasingly accessible, the temptation to deploy systems prematurely or overestimate their capabilities poses real risks to both users and markets. CryptoWhiz demonstrates that cutting-edge AI techniques can be implemented responsibly by being transparent about limitations, prioritizing risk management alongside profit seeking, providing clear reasoning for decisions to enable learning and oversight, and strongly communicating the educational purpose rather than encouraging real-money trading without extensive additional validation.

The cryptocurrency market, with its 24/7 trading, high volatility, and digital-native infrastructure, provides an ideal testing ground for AI trading systems before considering more traditional and regulated markets. The lessons learned from CryptoWhiz—about data pipeline engineering, model selection and training, risk management integration, and system architecture—transfer directly to applications in stock markets, forex trading, commodities, and other financial instruments. The project thus serves as both a specific cryptocurrency analysis tool and a general framework for AI-driven financial decision-making that can be adapted across domains.

Ultimately, CryptoWhiz succeeds in its primary educational mission by demystifying algorithmic trading, making advanced AI techniques accessible through open-source implementation, demonstrating the integration of multiple analytical perspectives into coherent decision-making, and maintaining intellectual honesty about what AI can and cannot achieve in financial markets. The project neither oversells AI as a magical solution to market complexity nor undersells it as entirely useless—instead, it occupies the realistic middle ground where AI provides valuable analytical tools that augment human judgment rather than replace it entirely.